

# K-Means Clustering Algorithm

September 12, 2023

## 0.1 K-Means Clustering Algorithm: Credit Card Customer Segmentation

In this project, I will play the role of a data scientist working for a credit card company. I have been given a dataset containing information about the company's clients and asked to help segment them into different groups in order to apply different business strategies for each type of customer.

The company expects to receive a group for each client and also an explanation of the characteristics of each group and what are the main points that make them different.

In a planning meeting with the Data Science coordinator, it was decided that I should use the K-means algorithm to segment the data.

In order to use the algorithm properly and achieve all the goals that the company has set for me, I will go through the following steps:

- Analyze the dataset;
- Prepare the data for modeling;
- Find an appropriate number of clusters;
- Segment the data;
- Interpret and explain the results.

I will start by importing the packages that I will use.

```
[38]: pip install kneed
```

```
Requirement already satisfied: kneed in /opt/conda/lib/python3.10/site-packages (0.8.5)
```

```
Requirement already satisfied: scipy>=1.0.0 in /opt/conda/lib/python3.10/site-packages (from kneed) (1.9.1)
```

```
Requirement already satisfied: numpy>=1.14.2 in /opt/conda/lib/python3.10/site-packages (from kneed) (1.21.6)
```

```
Note: you may need to restart the kernel to use updated packages.
```

```
[39]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from kneed import KneeLocator
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler
```

# 1 Exploring the Data

```
[40]: df = pd.read_csv('customer_segmentation.csv')
```

```
[41]: df.head()
```

```
[41]:
```

	customer_id	age	gender	dependent_count	education_level	marital_status	\
0	768805383	45	M	3	High School	Married	
1	818770008	49	F	5	Graduate	Single	
2	713982108	51	M	3	Graduate	Married	
3	769911858	40	F	4	High School	Unknown	
4	709106358	40	M	3	Uneducated	Married	

	estimated_income	months_on_book	total_relationship_count	\
0	69000	39	5	
1	24000	44	6	
2	93000	36	4	
3	37000	34	3	
4	65000	21	5	

	months_inactive_12_mon	credit_limit	total_trans_amount	\
0	1	12691.0	1144	
1	1	8256.0	1291	
2	1	3418.0	1887	
3	4	3313.0	1171	
4	1	4716.0	816	

	total_trans_count	avg_utilization_ratio
0	42	0.061
1	33	0.105
2	20	0.000
3	20	0.760
4	28	0.000

```
[42]: df.describe()
```

```
[42]:
```

	customer_id	age	dependent_count	estimated_income	\
count	1.012700e+04	10127.000000	10127.000000	10127.000000	
mean	7.391776e+08	46.325960	2.346203	62078.206774	
std	3.690378e+07	8.016814	1.298908	39372.861291	
min	7.080821e+08	26.000000	0.000000	20000.000000	
25%	7.130368e+08	41.000000	1.000000	32000.000000	
50%	7.179264e+08	46.000000	2.000000	50000.000000	
75%	7.731435e+08	52.000000	3.000000	80000.000000	
max	8.283431e+08	73.000000	5.000000	200000.000000	

	months_on_book	total_relationship_count	months_inactive_12_mon	\
--	----------------	--------------------------	------------------------	---

count	10127.000000	10127.000000	10127.000000
mean	35.928409	3.812580	2.341167
std	7.986416	1.554408	1.010622
min	13.000000	1.000000	0.000000
25%	31.000000	3.000000	2.000000
50%	36.000000	4.000000	2.000000
75%	40.000000	5.000000	3.000000
max	56.000000	6.000000	6.000000

	credit_limit	total_trans_amount	total_trans_count \
count	10127.000000	10127.000000	10127.000000
mean	8631.953698	4404.086304	64.858695
std	9088.776650	3397.129254	23.472570
min	1438.300000	510.000000	10.000000
25%	2555.000000	2155.500000	45.000000
50%	4549.000000	3899.000000	67.000000
75%	11067.500000	4741.000000	81.000000
max	34516.000000	18484.000000	139.000000

	avg_utilization_ratio
count	10127.000000
mean	0.274894
std	0.275691
min	0.000000
25%	0.023000
50%	0.176000
75%	0.503000
max	0.999000

Looking at the summary statistics here, the data makes sense for the variables is covers.

```
[43]: df.isnull().sum()
```

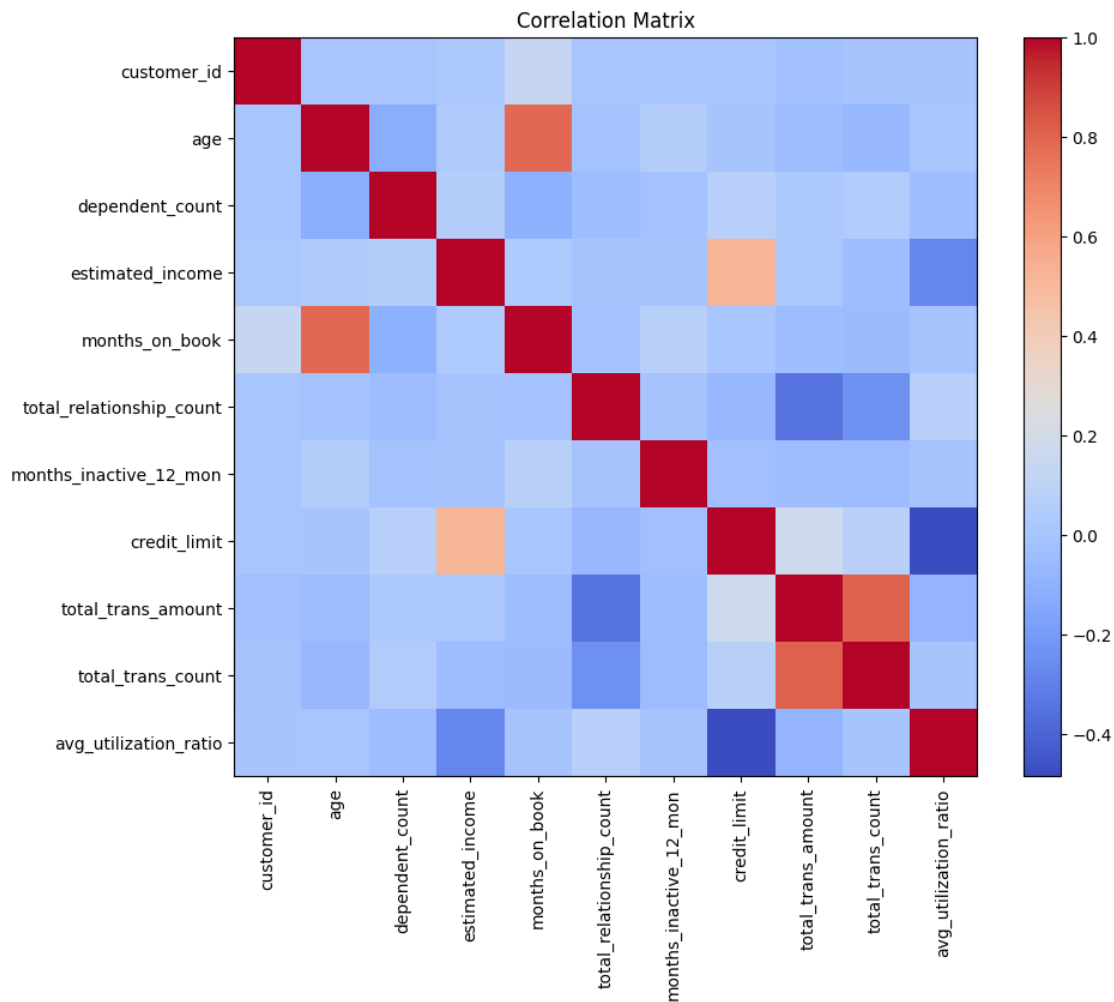
```
[43]: customer_id      0
      age              0
      gender           0
      dependent_count  0
      education_level  0
      marital_status   0
      estimated_income  0
      months_on_book   0
      total_relationship_count  0
      months_inactive_12_mon  0
      credit_limit      0
      total_trans_amount  0
      total_trans_count  0
      avg_utilization_ratio  0
```

dtype: int64

There aren't any missing values in the data.

```
[44]: correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
plt.imshow(correlation_matrix, cmap='coolwarm', interpolation='none',
           aspect='auto')
plt.colorbar()
plt.xticks(range(len(correlation_matrix)), correlation_matrix.columns,
           rotation=90)
plt.yticks(range(len(correlation_matrix)), correlation_matrix.columns)
plt.title('Correlation Matrix')
```

```
[44]: Text(0.5, 1.0, 'Correlation Matrix')
```



Most of the variables present weak correlations between each other, but there are some I can

highlight:

- Age is strongly correlated with how long the person has been a customer (months\_on\_book);
- Credit limit is positively correlated with the estimated income and negatively correlated with the average utilization ratio;
- The total number of transactions (total\_trans\_count) is strongly correlated with the total amount transitioned (total\_trans\_amount).

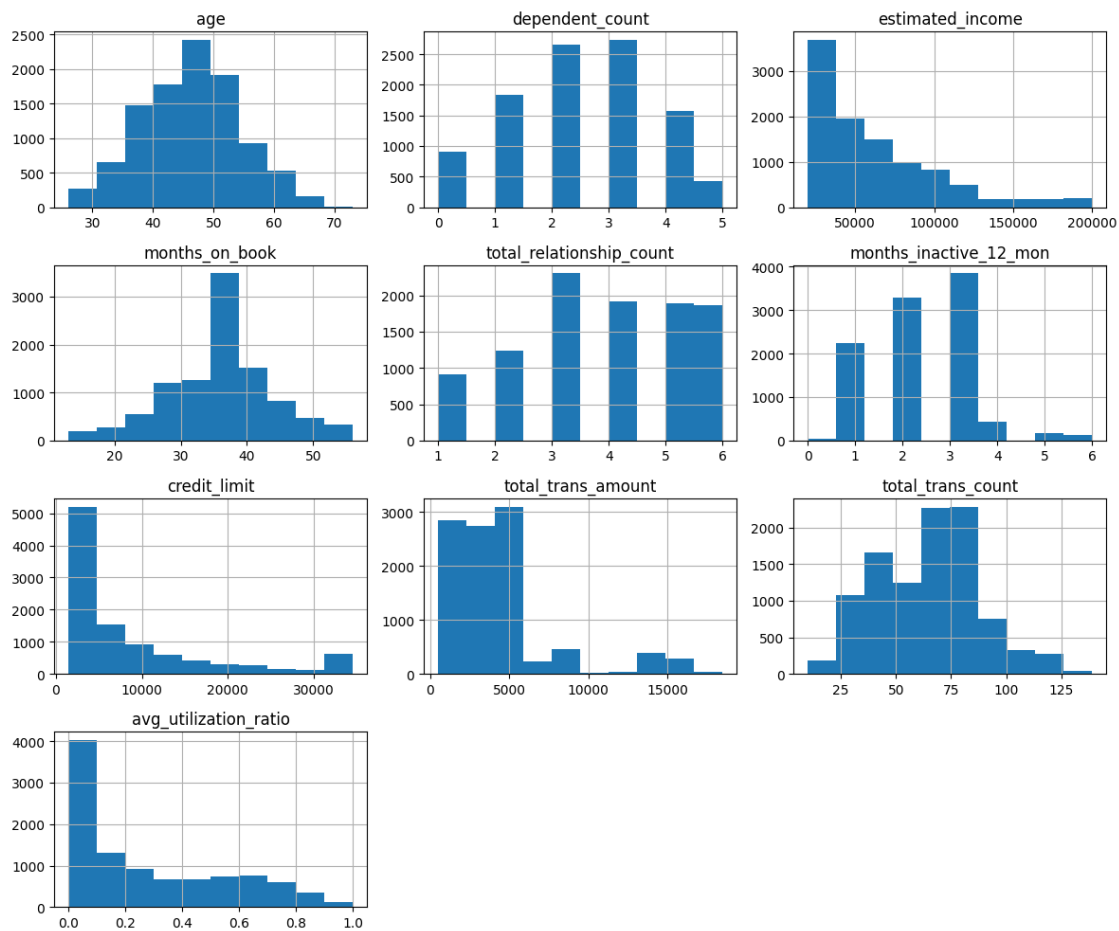
```
[45]: fig, ax = plt.subplots(figsize=(12, 10))

#Removing the customer's id before plotting the distributions
df.drop('customer_id', axis=1).hist(ax=ax)

plt.tight_layout()
plt.show()
```

/tmp/ipykernel\_64/579634717.py:4: UserWarning: To output multiple subplots, the figure containing the passed axes is being cleared

```
df.drop('customer_id', axis=1).hist(ax=ax)
```



The bar plots here show that most distributions in the dataframe are skewed.

## 1.1 Preprocessing: Feature Scaling using Z-Score Normalization

Now, I will scale all numerical columns using the `StandardScaler` class so they all have a mean of 0 and standard deviation of 1.

```
[46]: numerical_features = ['age', 'dependent_count', 'months_on_book',  
    ↪ 'total_relationship_count',  
    ↪ 'months_inactive_12_mon', 'credit_limit',  
    ↪ 'total_trans_amount',  
    ↪ 'total_trans_count', 'avg_utilization_ratio']  
  
# Extract the numerical data  
numerical_data = df[numerical_features]  
  
# Initialize the StandardScaler  
scaler = StandardScaler()  
  
# Fit and transform the data  
scaled_data = scaler.fit_transform(numerical_data)  
  
# Create a DataFrame with the scaled data and columns  
scaled_df = pd.DataFrame(data=scaled_data, columns=numerical_features)
```

I can see that all numerical features are now properly scaled. Now, I will add the gender column to `scaled_df` with the following mapping: 1 for M and 0 for F

```
[47]: scaled_df['gender'] = df['gender'].apply(lambda x: 1 if x == 'M' else 0)
```

```
[48]: scaled_df.head()
```

```
[48]:
```

	age	dependent_count	months_on_book	total_relationship_count	\
0	-0.165406	0.503368	0.384621	0.763943	
1	0.333570	2.043199	1.010715	1.407306	
2	0.583058	0.503368	0.008965	0.120579	
3	-0.789126	1.273283	-0.241473	-0.522785	
4	-0.789126	0.503368	-1.869317	0.763943	

	months_inactive_12_mon	credit_limit	total_trans_amount	\
0	-1.327136	0.446622	-0.959707	
1	-1.327136	-0.041367	-0.916433	
2	-1.327136	-0.573698	-0.740982	
3	1.641478	-0.585251	-0.951758	
4	-1.327136	-0.430877	-1.056263	

	total_trans_count	avg_utilization_ratio	gender
0	-0.973895	-0.775882	1

1	-1.357340	-0.616276	0
2	-1.911206	-0.997155	1
3	-1.911206	1.759686	0
4	-1.570365	-0.997155	1

Finally, I will map the `education_level` column in the `scaled_df` column in the following way:

- Uneducated - 0
- High School - 1
- College - 2
- Graduate - 3
- Post-Graduate - 4
- Doctorate - 5

```
[49]: education_mapping = {
      'Uneducated': 0,
      'High School': 1,
      'College': 2,
      'Graduate': 3,
      'Post-Graduate': 4,
      'Doctorate': 5
    }

scaled_df['education_level'] = df['education_level'].map(education_mapping)
```

```
[50]: scaled_df.head()
```

```
[50]:
```

	age	dependent_count	months_on_book	total_relationship_count	\
0	-0.165406	0.503368	0.384621	0.763943	
1	0.333570	2.043199	1.010715	1.407306	
2	0.583058	0.503368	0.008965	0.120579	
3	-0.789126	1.273283	-0.241473	-0.522785	
4	-0.789126	0.503368	-1.869317	0.763943	

	months_inactive_12_mon	credit_limit	total_trans_amount	\
0	-1.327136	0.446622	-0.959707	
1	-1.327136	-0.041367	-0.916433	
2	-1.327136	-0.573698	-0.740982	
3	1.641478	-0.585251	-0.951758	
4	-1.327136	-0.430877	-1.056263	

	total_trans_count	avg_utilization_ratio	gender	education_level
0	-0.973895	-0.775882	1	1
1	-1.357340	-0.616276	0	3
2	-1.911206	-0.997155	1	3
3	-1.911206	1.759686	0	1
4	-1.570365	-0.997155	1	0

Since the `marital_status` column in the original dataframe can't be separated in order of magnitude, I will use one-hot encoding on it.

```
[51]: marital_status_encoded = pd.get_dummies(df['marital_status'],
      ↪ prefix='marital_status')

scaled_df = pd.concat([scaled_df, marital_status_encoded], axis=1)
```

```
[52]: scaled_df.head()
```

```
[52]:
```

	age	dependent_count	months_on_book	total_relationship_count	\
0	-0.165406	0.503368	0.384621	0.763943	
1	0.333570	2.043199	1.010715	1.407306	
2	0.583058	0.503368	0.008965	0.120579	
3	-0.789126	1.273283	-0.241473	-0.522785	
4	-0.789126	0.503368	-1.869317	0.763943	

	months_inactive_12_mon	credit_limit	total_trans_amount	\
0	-1.327136	0.446622	-0.959707	
1	-1.327136	-0.041367	-0.916433	
2	-1.327136	-0.573698	-0.740982	
3	1.641478	-0.585251	-0.951758	
4	-1.327136	-0.430877	-1.056263	

	total_trans_count	avg_utilization_ratio	gender	education_level	\
0	-0.973895	-0.775882	1	1	
1	-1.357340	-0.616276	0	3	
2	-1.911206	-0.997155	1	3	
3	-1.911206	1.759686	0	1	
4	-1.570365	-0.997155	1	0	

	marital_status_Divorced	marital_status_Married	marital_status_Single	\
0	0	1	0	
1	0	0	1	
2	0	1	0	
3	0	0	0	
4	0	1	0	

	marital_status_Unknown
0	0
1	0
2	0
3	1
4	0

## 1.2 Choosing Number of Clusters (K) with Elbow Method

Now, I will choose the optimal K by looking at an Elbow Curve.



```
[53]: def plot_elbow_curve(df, max_clusters=10):
    inertias = []

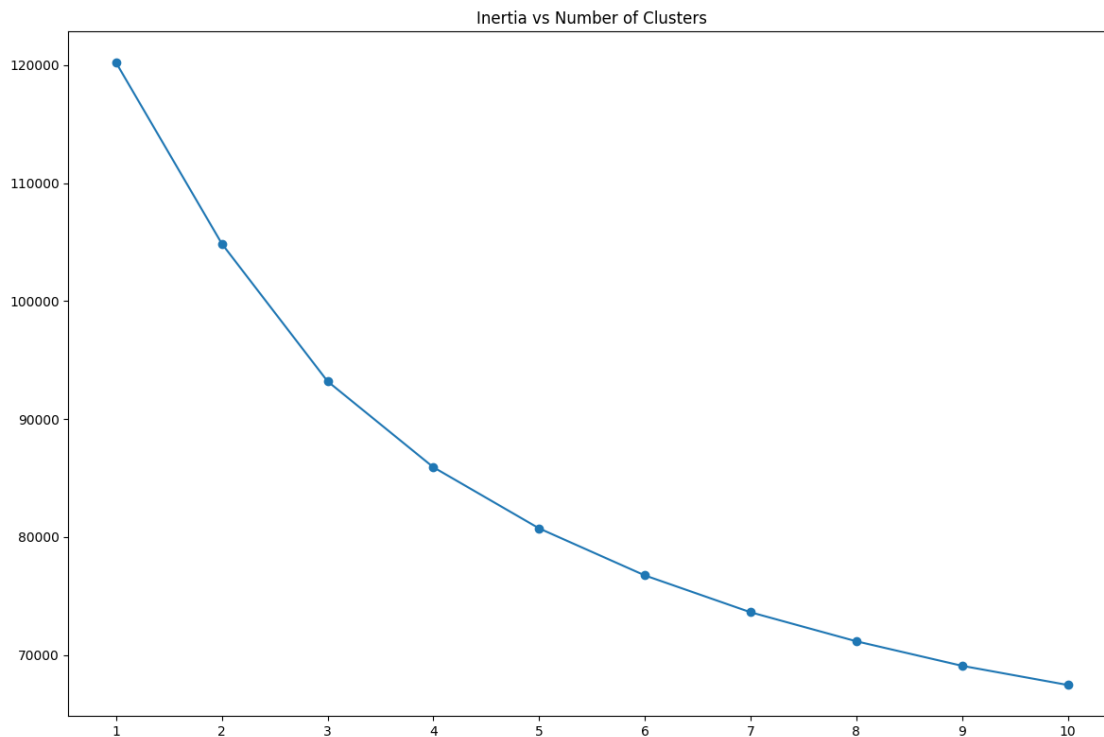
    for k in range(1, max_clusters+1):
        model = KMeans(n_clusters=k)
        cluster = model.fit_predict(df)
        inertias.append(model.inertia_)

    plt.figure(figsize=(12, 8))
    plt.plot(range(1, max_clusters+1), inertias, marker='o')
    plt.xticks(ticks=range(1, max_clusters+1), labels=range(1, max_clusters+1))
    plt.title('Inertia vs Number of Clusters')

    plt.tight_layout()
    plt.show()

    return inertias

inertias = plot_elbow_curve(scaled_df)
print(inertias)
```



```
[120209.61913696052, 104859.42237496469, 93190.27610220725, 85908.23031097917,
80722.76277888907, 76736.2446877103, 73610.56882778769, 71150.0175642037,
69071.21117731852, 67445.39805939124]
```

Looking at the graph, I will use the K value of 6.

### 1.3 Running K-Means Algorithm with K=6

```
[54]: model = KMeans(n_clusters=6)
clusters = model.fit_predict(scaled_df)

clusters_series = pd.Series(clusters)
df['Cluster'] = clusters_series + 1

print(df['Cluster'].value_counts())
```

```
6    2807
2    2381
4    1550
3    1288
5    1151
1     950
Name: Cluster, dtype: int64
```

After segmenting the data into 6 clusters, I can see that clusters 2, 4, and 3 have very similar number of occurrences, while clusters 5, 1, and 6 have very similar number of occurrences.

```
[55]: model.inertia_
```

```
[55]: 76736.14053004631
```

```
[56]: model.cluster_centers_
```

```
[56]: array([[ -0.13507309, -0.04043519, -0.10030536, -1.04627928, -0.15323135,
          0.62181075,  2.62704152,  1.79027379, -0.37989373,  0.61263158,
          2.01473684,  0.07684211,  0.44631579,  0.4         ,  0.07684211],
        [-0.01623693,  0.31347726, -0.01097579,  0.06678521,  0.01248237,
        -0.37535248, -0.20687004, -0.06862797,  0.17467158,  0.37184874,
          0.49579832,  0.07016807,  0.47058824,  0.37394958,  0.08529412],
        [-1.41090997, -0.66733314, -1.35727058,  0.33070788, -0.16947624,
        -0.34490732, -0.39578113, -0.32751519,  0.14885999,  0.4670287 ,
          2.15593483,  0.07680372,  0.44142746,  0.42668735,  0.05508146],
        [ 1.34858407, -0.9748692 ,  1.25024624,  0.18989628,  0.17249301,
        -0.35086044, -0.41509292, -0.4509529 ,  0.1655649 ,  0.42645161,
          2.27806452,  0.04129032,  0.51096774,  0.40451613,  0.04322581],
        [ 0.02767642,  0.23088507,  0.02595071,  0.08813115, -0.01578292,
          2.07295704, -0.25227897, -0.27502463, -0.84198649,  0.90695652,
          2.00608696,  0.09391304,  0.40695652,  0.41130435,  0.08782609],
        [-0.04861357,  0.4978844 , -0.03446964,  0.0046453 ,  0.03030727,
        -0.38919612, -0.19930984, -0.03561365,  0.1655837 ,  0.35470085,
          3.27029915,  0.08440171,  0.46794872,  0.36431624,  0.08333333]])
```

```
[57]: model.n_iter_
```

```
[57]: 29
```

When running the K-Means algorithm with K=4, I am able to produce the following results:

- Lowest SSE Value:85910.45
- The coordinates of the centroids
- Number of iterations required to converge: 7

```
[58]: model.labels_[:5]
```

```
[58]: array([1, 5, 5, 1, 1], dtype=int32)
```

## 1.4 Interpreting Results - Exploring Clusters for Numerical Variables

Now, I will interpret the results and summarize the characteristics of each cluster and differentiate them from each other based on the variables used for the segmentation.

Now, I will analyze the numerical variables and see how they behave in each cluster.

```
[59]: import matplotlib.pyplot as plt
import seaborn as sns

# Define the number of rows and columns in the subplot grid
num_rows = 2
num_cols = 5 # Increase the number of columns to 4 to accommodate 7 subplots

# Create a new figure
fig = plt.figure(figsize=(15, 10))

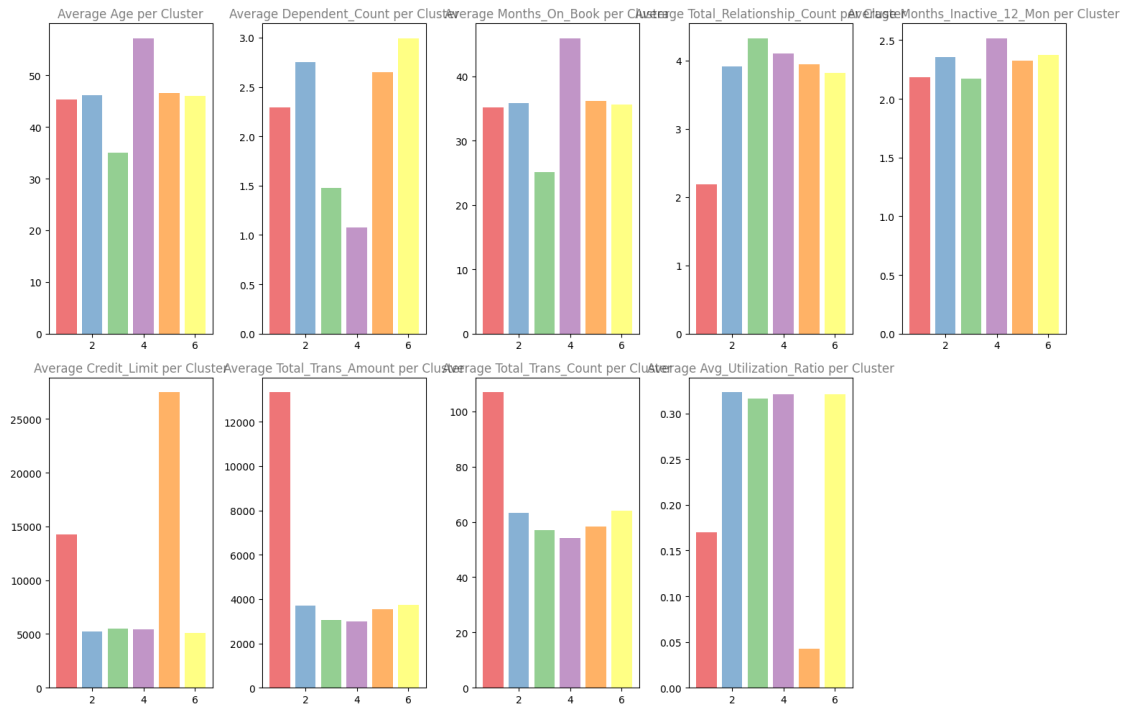
for i, column in enumerate(numerical_features):
    df_plot = df.groupby('Cluster')[column].mean()

    # Calculate the row and column index for the subplot
    row_index = i // num_cols
    col_index = i % num_cols

    # Add the subplot
    ax = fig.add_subplot(num_rows, num_cols, i + 1)

    # Plot the data in the subplot
    ax.bar(df_plot.index, df_plot, color=sns.color_palette('Set1'), alpha=0.6)
    ax.set_title(f'Average {column.title()} per Cluster', alpha=0.5)

plt.tight_layout()
plt.show()
```

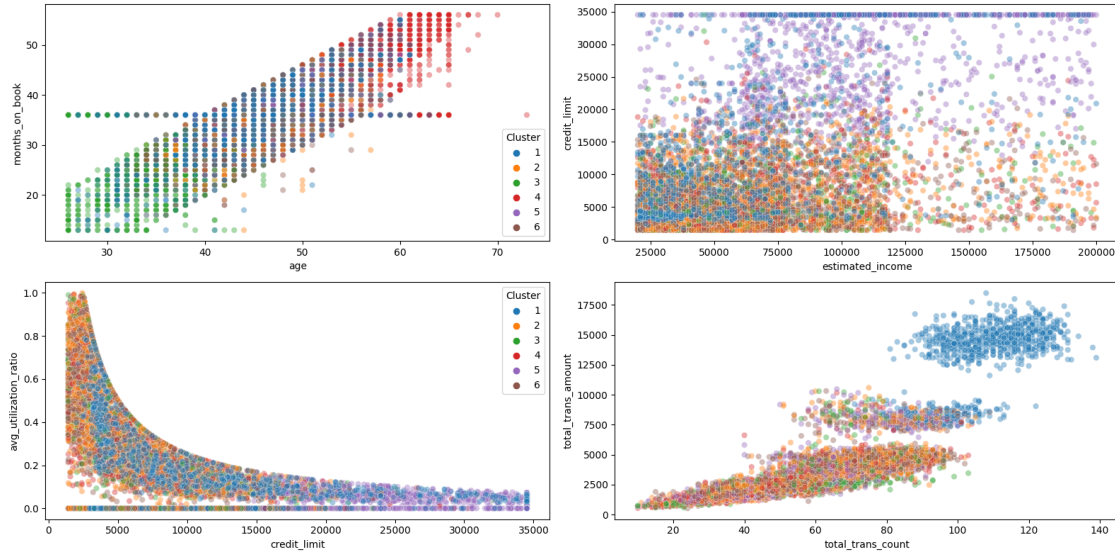


From looping the original dataframe, I found the following information:

- Cluster 4, on average, has the oldest clients while Cluster 3 has the youngest. These two clusters also have the lowest number of dependents, on average.
- Months on Book is the least for Cluster 3, on average which is expected.
- Cluster 1, on average, has the lowest relationship count but the highest transaction amounts.
- Cluster 5, on average, has the highest credit limit but also the lowest credit utilization rate.

```
[61]: fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(16, 8))
sns.scatterplot(x='age', y='months_on_book', hue='Cluster', data=df,
               palette='tab10', alpha=0.4, ax=ax1)
sns.scatterplot(x='estimated_income', y='credit_limit', hue='Cluster', data=df,
               palette='tab10', alpha=0.4, ax=ax2, legend=False)
sns.scatterplot(x='credit_limit', y='avg_utilization_ratio', hue='Cluster',
               data=df, palette='tab10', alpha=0.4, ax=ax3)
sns.scatterplot(x='total_trans_count', y='total_trans_amount', hue='Cluster',
               data=df, palette='tab10', alpha=0.4, ax=ax4, legend=False)

plt.tight_layout()
plt.show()
```



From these scatterplots, I can draw these conclusions:

- Cluster 1 has the highest amount of money transitioned.
- Cluster 2 has the lowest credit limit and estimated income and the highest utilization rate.
- Cluster 5 has the highest credit limit.
- Older clients are grouped in Cluster 5.

## 1.5 Interpreting Results - Exploring Clusters for Categorical Variables

Now, I want understand how the categorical column **Gender** impacts the cluster split.

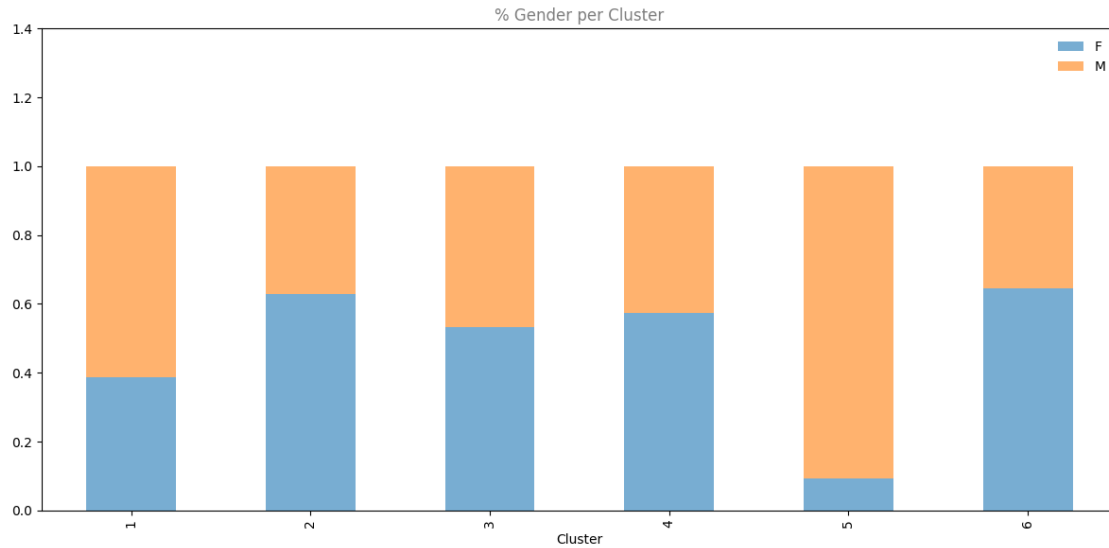
Here are the questions I want to answer: - Is the cluster with customers from only one gender? - Are any clusters equally divided between men and women?

```
[65]: plot_df = pd.crosstab(
    index=df['Cluster'], columns=df['gender'],
    values=df['gender'], aggfunc='size', normalize='index'
)

fig, ax = plt.subplots(figsize=(12,6))
plot_df.plot.bar(stacked=True, ax=ax, alpha=0.6)
ax.set_title(f'% Gender per Cluster', alpha=0.5)

ax.set_ylim(0, 1.4)
ax.legend(frameon=False)
ax.xaxis.grid(False)

plt.tight_layout()
plt.show()
```



For the categorical columns, I will plot the percentual distribution of each variable in each cluster.

```
[66]: cat_columns = df.select_dtypes(include=['object'])

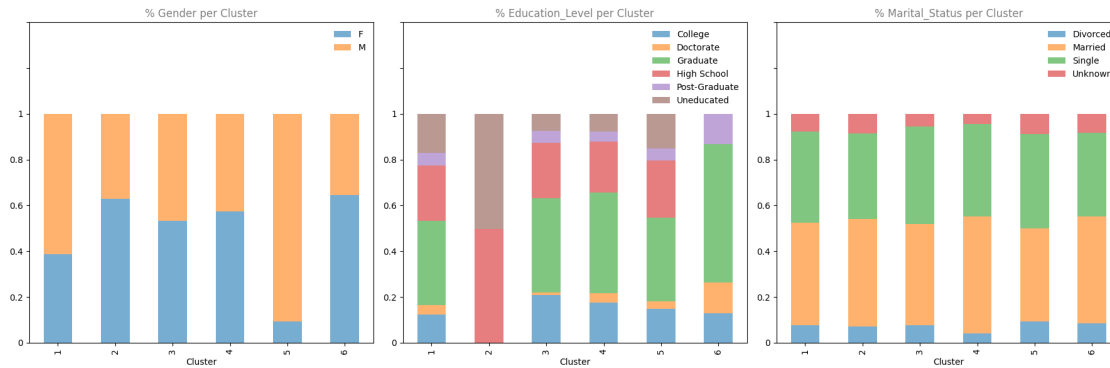
fig = plt.figure(figsize=(18, 6))
for i, col in enumerate(cat_columns):
    plot_df = pd.crosstab(index=df['Cluster'], columns=df[col], values=df[col],
        aggfunc='size', normalize='index')
    ax = fig.add_subplot(1, 3, i+1)
    plot_df.plot.bar(stacked=True, ax=ax, alpha=0.6)
    ax.set_title(f'% {col.title()} per Cluster', alpha=0.5)

    ax.set_ylim(0, 1.4)
    ax.legend(frameon=False)
    ax.xaxis.grid(False)

    labels = [0, 0.2, 0.4, 0.6, 0.8, 1]
    ax.set_yticklabels(labels)

plt.tight_layout()
plt.show()
```

```
/tmp/ipykernel_64/2465878031.py:15: UserWarning: FixedFormatter should only be
used together with FixedLocator
    ax.set_yticklabels(labels)
```



Considering the categorical variables, I notice:

- Cluster 5 is largely male, while cluster 6 is largely female.
- Cluster 2 is largely composed of high school graduates and those who are uneducated.
- Cluster has mostly graduates.
- Marital Status is evenly split across the clusters.

Now, I will evaluate the clustering performance of the model I ran above because the Elbow Method does not evaluate clustering performance using ground truth labels.

```
[ ]: from sklearn.cluster import DBSCAN
from sklearn.datasets import make_moons
from sklearn.metrics import adjusted_rand_score
```

```
[ ]: scaled_df, true_labels = make_moons(
    n_samples=250, noise=0.05, random_state=42
)
scaled_features = scaler.fit_transform(scaled_df)
```

```
[ ]: # Instantiate k-means and dbscan algorithms
kmeans = KMeans(n_clusters=4)
dbscan = DBSCAN(eps=0.3)

# Fit the algorithms to the features
kmeans.fit(scaled_features)
dbscan.fit(scaled_features)

# Compute the silhouette scores for each algorithm
kmeans_silhouette = silhouette_score(
    scaled_features, kmeans.labels_
).round(2)
dbscan_silhouette = silhouette_score(
    scaled_features, dbscan.labels_
).round(2)
```

```
[ ]: print(kmeans_silhouette)
      print(dbscan_silhouette)
```

The silhouette coefficient is higher for the K-Means algorithm. The DBSCAN algorithm appears to find more natural clusters according to the shape of the data.

The ARI metric uses true cluster assignments to measure the similarity between true and predicted labels. I will use this method to evaluate the clustering performance.

```
[ ]: ari_kmeans = adjusted_rand_score(true_labels, kmeans.labels_)
      ari_dbscan = adjusted_rand_score(true_labels, dbscan.labels_)

      round(ari_kmeans, 2)

[ ]: round(ari_dbscan, 2)
```

ARI shows that DBSCAN is the best choice compared to the K-Means because an ARI of 1 indicates perfectly labeled clusters.

## 1.6 Conclusion

As demanded by the company, I now have listed the most important characteristics of each cluster. Below are some suggestions and insights into each one of them.

### Cluster 1:

- Relationship Count: Lowest on average.
- Transaction Amounts: Highest on average.
- Suggestion: Cluster 1 seems to consist of clients with a lower number of relationships but higher transaction amounts. Consider tailoring marketing strategies to these clients to maximize transaction revenue.

### Cluster 2:

- Credit Limit: Lowest on average.
- Estimated Income: Lowest on average.
- Credit Utilization Rate: Highest on average.
- Suggestion: Cluster 2 appears to have clients with lower credit limits, lower income, and higher credit utilization rates. Offer financial education or credit management services to help clients in this cluster improve their financial health.

### Cluster 3:

- Average Age: Youngest on average.
- Average Number of Dependents: Lowest on average.
- Months on Book: Least on average.
- Suggestion: Cluster 3 consists of younger clients who have been with the bank for a shorter time. Focus on building long-term relationships with these clients to maximize their value over time.

### Cluster 4:

- Average Age: Oldest on average.



- Average Number of Dependents: Lowest on average.
- Suggestion: Cluster 4 has the oldest clients with the fewest dependents. Consider offering retirement planning or wealth management services to cater to the needs of this demographic.

#### **Cluster 5:**

- Credit Limit: Highest on average.
- Credit Utilization Rate: Lowest on average.
- Suggestion: Cluster 5 appears to consist of clients with high credit limits and low credit utilization rates. Encourage these clients to make use of their credit lines responsibly and offer premium financial products.

#### **Cluster 6:**

- Gender: Largely female.
- Suggestion: Cluster 6 is predominantly female. Consider tailoring marketing campaigns or services to the specific needs and preferences of female clients.

#### **Education and Marital Status:**

- Cluster 2: Largely composed of high school graduates and those who are uneducated.
- Cluster 4: Mostly graduates.
- Marital Status is evenly split across the clusters.

#### **Additional Insights:**

- Cluster 5 has the highest credit limit and consists of older clients. These clients might be financially stable and could be targeted for premium services.
- Cluster 2 has lower credit limits and higher credit utilization rates. Consider offering credit counseling or credit limit increase options.
- Marital status doesn't appear to vary significantly across clusters, so marital status may not be a strong indicator of behavior or needs in this context.

It's important to conduct further analysis, possibly using statistical techniques or machine learning models, to identify patterns and trends more precisely and to validate these insights. Additionally, consider conducting surveys or interviews with clients in each cluster to gather more specific information about their financial needs and preferences.

[ ]: